


Where naïve proposals go — every rejection routed to a deterministic remedy

Mini-benchmark · 24 seeded naïve proposals · adversarial stress corpus (pathological by design)



Driven by autocarto.benchmark.build_report() — no fabricated flows; counts sum to the real scenarios. The corpus is deliberately adversarial, so the rejection share is high by design, not a natural rate. First-pass verdicts (not the full iterate-to-convergence loop). Of the 6 Gate-3b approvals, 1 is a known false-approval — two independent autocorrelated fields correlating by chance (the documented free-permutation null-model limitation, research task R-2).