One rejection, end to end — the machine-readable audit trail

Gate 2 · zero-inflated variable (e.g. asthma hospitalisations) · Propose → Verify → Execute

1 PROPOSAL · simulated LLM (fed directly to Gate 2)

```
"proposed_method": "jenks",
"proposed_breaks":  [0.00, 1.53, 3.95, 68.88]

# naïve Fisher-Jenks on the raw, zero-inflated values
# Tier 1 is SIMULATED in V1 — the harness feeds this straight to Gate 2
```

▼ *profile + verify*

2 GATE 2 VERDICT · verbatim from the emitted JSON trace

```
"gate":              "G2",
"diagnosis":         "zero_inflated",
"passed":            false,       # 49.8% of tracts are zero
"prescribed_method": "manual_break_at_zero_then_fisher_jenks",
"prescribed_breaks": [0.0, 7.86, 20.61, 68.88],
"instruction": "Mandate explicit break at 0, then Fisher-Jenks
   on the non-zero tail.
   DO NOT propose alternative methods.
   Use these exact breaks."
```

▼ *MANDATE (prescribed breaks + code)*

3 EXECUTED CODE · the mandate, transcribed

```
- breaks = [0.00, 1.53, 3.95, 68.88]        # LLM proposal — REJECTED

+ # MANDATED CLASSIFICATION — DO NOT MODIFY

+ breaks = [0.0, 7.86, 20.61, 68.88]        # prescribed by Gate 2

+ classified = np.digitize(values, bins=breaks, right=True)
```

The LLM is reduced to a code-assembler: it may only splice in the prescribed constants.

Verbatim from output/traces/gate2_classification_trace.json (case: zero_inflated); regenerate with `autocarto demo`.
Break values rounded for display — the trace stores full float precision. Re-running reproduces the trace byte-for-byte.